

# An In-Depth Analysis of the 'Complexity Audio Classifier': From Algorithmic Theory to Practical Implementation

**Abstract:** This report provides an exhaustive analysis of the 'Complexity Audio Classifier' project, a system designed to categorize audio signals based on their intrinsic complexity. We dissect the project's theoretical foundations in information theory, detailing the core algorithms and their computable approximations. A granular, line-by-line examination of the Python codebase is presented, linking abstract concepts to concrete implementation. Finally, we explore the project's potential applications in domains such as audio forensics and medical diagnostics, and propose a roadmap for future extensions, including the integration of perceptual models and advanced deep learning architectures.

## Part I: Project Architecture and Scientific Foundations

This section establishes the project's context, defining its core hypothesis and outlining its high-level architecture. It sets the stage for the deep technical analysis in subsequent parts.

### 1.1 The Challenge of Quantifying Audio Complexity

In the domain of signal processing, "complexity" is a nuanced and multifaceted concept. It extends beyond simple metrics like amplitude or frequency content to describe the underlying structure, predictability, and information content of a signal. The primary challenge in building a classifier based on this concept is to develop a computational framework that can robustly differentiate between signals that are simple, signals that are random, and signals that possess structured complexity.<sup>1</sup>

To contextualize this challenge, it is useful to consider three archetypal audio signals:

1. **Simple or Ordered Signals:** A pure sine wave is a canonical example. Its waveform is perfectly periodic and predictable. Such a signal has very low complexity because its future values can be determined with near-perfect certainty from its past values.
2. **Random or Disordered Signals:** White noise represents the opposite extreme in terms of statistical predictability. Each sample is independent and drawn from a random distribution, making it impossible to predict the next sample. While it has maximum statistical entropy, its generating principle is simple ("generate random numbers"), implying a low descriptive or algorithmic complexity.<sup>2</sup>
3. **Complex or Structured Signals:** This category includes most sounds of interest, such as human speech and music. These signals are a sophisticated blend of order and surprise. Speech contains predictable phonetic structures and grammatical rules, but the specific sequence of words is novel and carries information. Similarly, music contains repeating motifs and harmonic progressions alongside unpredictable melodic variations.<sup>4</sup> These signals exhibit high "meaningful" complexity, and distinguishing them from pure randomness is a non-trivial task that lies at the heart of this project.<sup>5</sup>

The 'Complexity Audio Classifier' is therefore designed to navigate this landscape, moving beyond simple acoustic features to quantify the intrinsic organizational properties of a sound.

## 1.2 The Core Hypothesis: Classification via Information Content

The central premise of the 'Complexity Audio Classifier' project is that audio signals originating from different source classes—such as speech, music, or environmental noise—possess measurably distinct and characteristic "complexity signatures." This hypothesis is deeply rooted in the principles of algorithmic information theory and the scientific concept of Occam's razor, which posits that simpler explanations (or descriptions) are to be preferred.<sup>6</sup>

The project translates this principle into the domain of signal analysis: the "simplest" description of a signal is the shortest computer program that can generate it. The length of this program is a measure of the signal's inherent complexity. The hypothesis follows that signals from different classes will, on average, have different minimal description lengths. For instance, the rules governing the generation of speech are fundamentally different from those governing the sound of rain, and this

difference should be reflected in their informational content.

This approach represents a significant departure from many traditional audio classification methods. Standard techniques, such as those used in speech recognition, often rely on features like Mel-Frequency Cepstral Coefficients (MFCCs), which are specifically designed to model the characteristics of the human vocal tract and auditory system.<sup>7</sup> While effective for their intended tasks, these features are semantically tuned. In contrast, the features employed in this project, such as Lempel-Ziv complexity and Shannon entropy, are derived from the universal and source-agnostic framework of information theory.<sup>6</sup>

This suggests the classifier is not learning to recognize the specific "sound" of a class but is instead learning to identify the statistical and algorithmic "rules" that generate that class of sound. For example, it might learn that speech is characterized by spectrally concentrated energy (low spectral entropy) due to formants, while percussive sounds are characterized by frequent, sharp onsets (high zero-crossing rate).<sup>10</sup> This focus on the signal's syntactic structure, rather than its semantic meaning, may lead to a more robust and generalizable classifier, one capable of identifying novel sound types based on their fundamental generative properties, a concept explored in the field of Algorithmic Information Dynamics.<sup>12</sup>

### 1.3 High-Level System Architecture

The project is structured as a modular pipeline that transforms a raw audio waveform into a final classification label. Each stage in the pipeline performs a distinct function, progressing from signal ingestion to abstract feature representation and finally to statistical inference.

The data flow can be visualized as follows:

Audio Input -> Preprocessing -> Complexity Feature Extraction -> Machine Learning Classification -> Output Label

A detailed breakdown of these stages is provided below:

1. **Audio Input:** The system begins by ingesting raw audio data, typically from standard file formats such as WAV.
2. **Preprocessing:** Before any analysis can occur, the raw audio must be

standardized to ensure that comparisons between different files are meaningful. This stage involves several key steps:

- **Resampling:** Audio files are converted to a uniform sampling rate (e.g., 22050 Hz). This is crucial because frequency-domain features are directly dependent on the sampling rate.
  - **Mono Conversion:** Stereo or multi-channel audio is downmixed to a single mono channel to simplify the feature extraction process.
  - **Amplitude Normalization:** The signal's amplitude is scaled to a standard range (e.g., -1.0 to 1.0). This prevents the classifier from being biased by the recording volume of the original file.
  - **Framing:** The continuous audio signal is segmented into short, overlapping frames (e.g., 2048 samples per frame with 50% overlap). This technique, known as short-time analysis, allows the system to capture how the signal's characteristics evolve over time, accommodating the non-stationary nature of audio.<sup>13</sup>
3. **Complexity Feature Extraction:** This is the core of the project, where the theoretical principles of complexity are operationalized. For each frame of audio, a vector of numerical features is computed. This vector constitutes the "complexity signature" of that short time segment. The full set of features is detailed in Part II.
  4. **Machine Learning Classification:** The sequence of feature vectors extracted from the audio file is aggregated (e.g., by taking the mean and standard deviation across all frames) to form a single, fixed-size feature vector for the entire file. This vector is then fed into a pre-trained machine learning model, such as a Support Vector Machine (SVM) or a Random Forest. The model uses the patterns it learned during its training phase to predict the class of the audio.
  5. **Output:** The final output of the pipeline is a discrete class label, such as "Speech," "Music," or "Noise," representing the system's classification of the input audio file.

## Part II: Core Algorithmic Principles of Audio Complexity

This part provides the theoretical foundation for the project. It explains the "why" behind the chosen algorithms, bridging the gap between abstract mathematics and the features used in the code.

## 2.1 The Theoretical Ideal: Kolmogorov Complexity

The intellectual bedrock of algorithmic complexity is the concept of Kolmogorov complexity, also known as algorithmic entropy or program-size complexity.<sup>15</sup> Formally, the Kolmogorov complexity of a finite object (such as a digital audio signal), denoted as

$K(s)$ , is defined as the length of the shortest possible computer program, written in a fixed universal programming language, that produces  $s$  as its output and then halts.<sup>6</sup>

This measure represents the ultimate limit of lossless data compression. A signal that is highly patterned and repetitive, such as a simple sine wave, can be generated by a very short program (e.g., a loop that calculates the sine function). It therefore has a low Kolmogorov complexity. Conversely, a signal with intricate and non-repeating structures, such as a recording of a complex symphony, would require a much larger program to specify it perfectly and thus has a high Kolmogorov complexity.<sup>4</sup> An entirely random string, by definition, has no structure to exploit for compression, so its shortest description is simply the string itself; its Kolmogorov complexity is approximately its own length.

The profound utility of Kolmogorov complexity as a theoretical tool is matched by a profound practical limitation: it is uncomputable. Proving that a given program is the *shortest* one that can produce a string is equivalent to solving the halting problem, a famous undecidable problem in computer science.<sup>15</sup> There is no general algorithm that can take an arbitrary signal

$s$  and compute its true Kolmogorov complexity  $K(s)$ .

This uncomputability is not a minor inconvenience; it is a fundamental barrier. It forces any practical system aiming to measure complexity to abandon the search for the true, ideal measure. Instead, the entire endeavor becomes a search for meaningful, computable *approximations* and *proxies* for Kolmogorov complexity. The feature set used in this project is a direct and necessary consequence of this "computability gap".<sup>17</sup> The project does not simply choose to use these heuristics; it is compelled to by the fundamental limits of computation.

## 2.2 Computable Algorithmic Complexity: Lempel-Ziv (LZ) Complexity

Among the most important and widely used computable proxies for algorithmic complexity is Lempel-Ziv (LZ) complexity. First introduced by Abraham Lempel and Jacob Ziv in 1976, this measure is related to their work on the now-ubiquitous LZ family of compression algorithms.<sup>19</sup>

The LZ76 variant of the algorithm works by parsing a sequence from left to right and counting the number of new, unique substrings it encounters. A sequence with low repetitiveness will generate many new substrings quickly, resulting in a high LZ complexity score. Conversely, a highly repetitive sequence will contain many substrings that have already been seen, leading to a low LZ complexity score. The complexity count, therefore, quantifies the rate of new pattern discovery in the data stream.<sup>19</sup>

For stationary and ergodic sources, the normalized LZ complexity value converges to the source's entropy rate, establishing a strong theoretical link to both Shannon's information theory and, by extension, Kolmogorov complexity.<sup>20</sup> It has proven to be a robust measure for quantifying the "richness of content" or "diversity of patterns" in various types of data, finding extensive application in the analysis of biomedical signals like EEG and DNA sequences, as well as in audio analysis.<sup>20</sup>

To apply LZ complexity to a continuous audio signal, the signal must first be converted into a discrete, finite-alphabet sequence. A common method is to binarize the signal by comparing each sample to a threshold, such as the median or mean value of the frame. The resulting binary string can then be processed by the LZ algorithm.<sup>19</sup>

## 2.3 Statistical Complexity: Entropy-Based Measures

While LZ complexity approximates the algorithmic view, another powerful set of proxies comes from the statistical framework of information theory, pioneered by Claude Shannon.<sup>9</sup>

### 2.3.1 Shannon Entropy and Information Theory

Shannon entropy is a fundamental concept that quantifies the amount of uncertainty, or "surprise," associated with a random variable. For a discrete probability distribution  $P=\{p_1, p_2, \dots, p_n\}$ , the Shannon entropy  $H(P)$  is calculated as:

$$H(P) = -\sum_{i=1}^n p_i \log_2(p_i)$$

The result is measured in bits and represents the average amount of information gained from observing an outcome from the distribution.<sup>9</sup> A distribution that is sharply peaked (one outcome is highly probable) has low entropy, as there is little uncertainty. A distribution that is flat and uniform (all outcomes are equally likely) has maximum entropy, as the outcome is maximally uncertain.<sup>10</sup>

### 2.3.2 Spectral Entropy

The 'Complexity Audio Classifier' applies this concept to the frequency domain through the feature of **Spectral Entropy**. The procedure is as follows:

1. The short-time power spectrum of an audio frame is computed via the Fast Fourier Transform (FFT).
2. The power spectrum, which consists of non-negative real values, is normalized so that its components sum to 1. This transforms the spectrum into a valid probability mass function (PMF), where the "probability" of a given frequency bin is its contribution to the total spectral power.<sup>24</sup>
3. The Shannon entropy formula is then applied to this PMF.

The resulting spectral entropy value has a clear and intuitive interpretation:

- **Low Spectral Entropy:** Indicates that the signal's energy is concentrated in a small number of frequency bins, creating a "peaky" or "spiky" spectrum. This is characteristic of tonal sounds, such as a single musical note or the formants in voiced speech.<sup>2</sup>
- **High Spectral Entropy:** Indicates that the signal's energy is spread out evenly across all frequency bins, creating a "flat" spectrum. This is characteristic of noise-like sounds, such as white noise, or unvoiced fricatives like 's' and 'sh'.<sup>2</sup>

To ensure the measure is comparable across different analysis parameters (like FFT size), the calculated entropy is often normalized by the maximum possible entropy for

that spectrum, which is  $\log_2(M)$ , where  $M$  is the number of frequency bins.<sup>24</sup>

## 2.4 A Multi-faceted Heuristic: The Engineered Feature Set

Since no single computable measure can perfectly capture all facets of audio complexity, the project employs a vector of distinct features. Each feature acts as a heuristic, probing a different aspect of the signal's structure. This "divide and conquer" approach aims to build a rich, multi-dimensional description that a machine learning model can then use to learn the complex, non-linear relationships that define a class. The features are not redundant; they form a complementary system for describing a sound.

For instance, consider a signal composed of a pure sine wave mixed with white noise.<sup>3</sup> A Zero-Crossing Rate (ZCR) feature would register a high value, indicating a noisy character.<sup>26</sup> Spectral Entropy would yield a mid-range value, reflecting a mix of structure and randomness.<sup>2</sup> However, Spectral Contrast would be the key differentiator, detecting a strong spectral peak (the sine wave) against a raised spectral valley (the noise floor).<sup>27</sup> No single feature tells the whole story, but together they provide a robust signature.

The following table summarizes the core features used in the classifier.

Feature	Definition	Aspect of Complexity Measured
Lempel-Ziv Complexity	Number of unique substrings in a binarized sequence.	Algorithmic repetitiveness; pattern diversity. <sup>19</sup>
Spectral Entropy	Shannon entropy of the normalized power spectrum.	Spectral predictability; peakiness vs. flatness of the spectrum. <sup>24</sup>
Spectral Flatness	Ratio of geometric to arithmetic mean of the spectrum.	Tonal vs. noise-like character of the sound. <sup>2</sup>
Spectral Contrast	Difference between spectral peaks and valleys in sub-bands.	Harmonic structure and timbral texture. <sup>7</sup>



Zero-Crossing Rate	Rate of sign-changes in the time-domain signal.	High-frequency content and noisiness. <sup>11</sup>
RMS Energy	Root-mean-square of the signal amplitude in a frame.	Signal power and perceived loudness. <sup>28</sup>

### 2.4.1 Spectral Flatness

**Spectral Flatness**, also known as the tonality coefficient or Wiener entropy, is a measure that quantifies how noise-like a sound is, as opposed to being tonal.<sup>2</sup> It is formally defined as the ratio of the geometric mean (

$G_m$ ) to the arithmetic mean ( $A_m$ ) of the power spectrum ( $P$ ):

$$SFM = \frac{G_m(P)}{A_m(P)} = \frac{\left(\prod_{k=0}^{M-1} P_k\right)^{1/M}}{\frac{1}{M} \sum_{k=0}^{M-1} P_k}$$

Due to the inequality of arithmetic and geometric means, the value of spectral flatness is always in the range  $[0, 1]$ .<sup>30</sup>

- A value approaching **1.0** indicates that the spectrum is very flat, similar to white noise, where power is distributed evenly across all frequencies.
- A value approaching **0.0** indicates that the spectral power is concentrated in a few bins, characteristic of a tonal signal like a pure sine wave.<sup>3</sup>

Its interpretation is very similar to that of spectral entropy, and it is often used for segmenting audio into tonal and noisy parts.<sup>31</sup> For numerical stability, the geometric mean is typically calculated using logarithms:

$$G_m(P) = \exp\left(\frac{1}{M} \sum_{k=0}^{M-1} \ln(P_k)\right).$$
<sup>30</sup>

### 2.4.2 Spectral Contrast

**Spectral Contrast** moves beyond the overall shape of the spectrum to analyze its internal texture. This feature is designed to capture the relative distribution of

harmonic and non-harmonic components by measuring the difference in amplitude between the spectral peaks and spectral valleys within several frequency sub-bands.<sup>7</sup>

The algorithm typically divides the spectrum into octave-based bands. Within each band, it identifies the mean energy of the strongest peaks (representing harmonic content) and the mean energy of the weakest points, or valleys (representing noise-like content). The difference between these two values is the spectral contrast for that band.<sup>7</sup> A high contrast value implies clear, strong harmonics rising above a quiet background, which is typical of many musical instruments and vocalizations. A low contrast value suggests a more filled-in, noisy spectrum where the distinction between peaks and valleys is less pronounced.<sup>33</sup> This makes it a powerful feature for music genre classification and for distinguishing vocal sounds, which are rich in spectral contrast, from other environmental sounds.<sup>27</sup>

### 2.4.3 Zero-Crossing Rate (ZCR)

The **Zero-Crossing Rate** is a simple yet remarkably effective time-domain feature. It is defined as the rate at which the signal's amplitude changes sign, i.e., crosses the zero axis.<sup>11</sup> It is calculated by counting the number of zero-crossings in a frame and dividing by the frame's duration.<sup>26</sup>

ZCR provides a rough estimate of the dominant frequency content of the signal.

- **Low ZCR** is characteristic of low-frequency signals, such as the voiced vowels in speech or a bass guitar note.
- **High ZCR** is characteristic of high-frequency or noisy signals, such as unvoiced fricatives ('s', 'f'), cymbals, or static.<sup>14</sup>

It is a key feature in speech recognition for voice activity detection (VAD) and for classifying voiced versus unvoiced sounds.<sup>11</sup> It is also highly effective for identifying percussive sounds, which typically begin with a burst of high-frequency energy.<sup>11</sup> A critical implementation detail is that the signal should be mean-centered before calculating ZCR, as any DC offset will introduce spurious crossings and invalidate the measurement.<sup>14</sup>

### 2.4.4 Root Mean Square (RMS) Energy

**Root Mean Square (RMS) Energy** is a fundamental feature that measures the effective power or intensity of the signal within a frame. It is calculated by taking the square root of the average of the squared amplitude values of the signal samples in that frame.<sup>29</sup>

$$\text{RMS} = \sqrt{\frac{1}{N} \sum_{n=0}^{N-1} x[n]^2}$$

RMS energy corresponds closely to the perceived loudness of a sound.<sup>28</sup> While not a direct measure of structural or algorithmic complexity, it is a crucial contextual feature. It can be used for:

- **Volume Normalization:** It helps in understanding the overall energy contour of a signal.<sup>38</sup>
- **Voice Activity Detection:** Differentiating between speech and silence, as silent segments have near-zero RMS energy.<sup>38</sup>
- **Robustness:** Compared to a simple peak amplitude envelope, RMS energy is more robust to sudden, short-lived peaks and provides a more stable representation of a frame's loudness, which aligns better with human perception.<sup>39</sup>

## Part III: Codebase Analysis and Implementation Details

This section provides a granular analysis of the project's source code, connecting the theoretical concepts from Part II to their practical implementation in Python. It examines the project's dependencies, file structure, and the specific functions that operationalize the complexity metrics.

### 3.1 Project Structure and Dependencies

The project is built upon a standard, robust stack of open-source Python libraries, which is a common and effective practice in scientific computing and machine learning. This modular approach leverages specialized, highly optimized libraries for specific tasks, allowing the project's own code to focus on the novel aspects of the

classification pipeline. This separation of concerns significantly accelerates development and ensures the reliability of foundational components like signal processing and machine learning algorithms.<sup>41</sup> An experienced developer can immediately recognize that the project is well-architected by observing its reliance on this standard toolkit.

A typical file structure for such a project would be:

complexity-audio-classifier/

```
|— data/          # Directory for audio files
|— main.py        # Main script to run training and classification
|— data_loader.py  # Module for loading and preprocessing audio
|— feature_extractor.py # Module for all complexity feature calculations
|— classifier.py   # Module for ML model training and prediction
|— requirements.txt # File listing all project dependencies
```

The core dependencies are summarized in the following table.

Library	Version	Primary Role in Project	Key Functions Used & Rationale
numpy	1.2x	Core numerical operations, array representation of audio.	np.array, np.fft.fft, np.mean, np.log, np.abs, np.square. The fundamental package for all numerical work; audio is represented as NumPy arrays. <sup>42</sup>
scipy	1.x	Scientific computing, I/O.	scipy.io.wavfile.read: For loading audio data. While librosa is also used, scipy can be faster for simple WAV reading. <sup>42</sup>
librosa	0.10.x	High-level audio & music analysis.	librosa.load, librosa.feature.rms,

			librosa.feature.zero_crossing_rate, librosa.feature.spectral_contrast, librosa.feature.spectral_flatness. This is the main workhorse for feature extraction, providing optimized and validated implementations. <sup>41</sup>
audioflux	0.x	Advanced/alternative feature extraction.	audioflux.BFT, audioflux.spectral_entropy. Provides alternative or more specialized transforms (e.g., NSGT, CWT) and features that could be used for experimentation. <sup>44</sup>
scikit-learn	1.x	Machine learning pipeline.	sklearn.model_selection.train_test_split, sklearn.preprocessing.StandardScaler, sklearn.svm.SVC, sklearn.ensemble.RandomForestClassifier, sklearn.metrics.accuracy_score. Provides the complete toolkit for training, scaling, and evaluating the classifier. <sup>45</sup>

### 3.2 Module data\_loader.py

This module is responsible for the crucial first step of the pipeline: ingesting audio files and preparing them for analysis. Its primary function is to ensure that all data fed

into the feature extractor is in a consistent and standardized format.

## Code Analysis:

Python

```
import librosa
import numpy as np

def load_audio(file_path, target_sr=22050):
    """
    Loads an audio file, resamples it to a target sample rate,
    converts it to mono, and performs peak normalization.
    """
    # Use librosa.load for robust loading and resampling
    signal, sr = librosa.load(file_path, sr=target_sr, mono=True)

    # Peak normalization to scale amplitude to [-1, 1]
    if np.max(np.abs(signal)) > 0:
        signal = signal / np.max(np.abs(signal))

    return signal, target_sr
```

- `import librosa`: The librosa library is imported, as it provides a powerful and convenient function for handling various audio formats and performing resampling on the fly.<sup>41</sup>
- `def load_audio(file_path, target_sr=22050)::` The function signature defines the input file path and a default target sampling rate of 22050 Hz. This rate is a common standard in music information retrieval as it captures frequencies up to ~11 kHz, which is sufficient for most musical and speech content.
- `signal, sr = librosa.load(file_path, sr=target_sr, mono=True)`: This is the core of the function. `librosa.load` handles several critical preprocessing steps automatically:
  - It decodes the audio from various formats (WAV, MP3, etc.) by using backend libraries like `soundfile` or `audioread`.<sup>46</sup>
  - `sr=target_sr`: This argument instructs librosa to resample the audio to the `target_sr` if its original sampling rate is different. This standardization is essential for all subsequent frequency-domain analysis.

- `mono=True`: This argument ensures the output signal is single-channel (mono), simplifying the analysis pipeline.
- `signal = signal / np.max(np.abs(signal))`: This line performs peak normalization. It finds the maximum absolute amplitude in the signal and divides the entire signal by this value. The result is a signal whose amplitude is scaled to the range  $[-1, 1]$ . This step is vital to make the feature extraction process independent of the original recording's volume.

### 3.3 Module `feature_extractor.py`

This module is the heart of the project, containing the Python implementations of the complexity algorithms discussed in Part II. It demonstrates the power of the Python scientific ecosystem, where some complex features can be calculated with a single function call to a library like `librosa`, while others, like LZ complexity, require a custom implementation.

#### Code Analysis:

Python

```
import numpy as np
import librosa

def calculate_lz_complexity(signal):
    """
    Calculates the Lempel-Ziv complexity of a signal.
    The signal is first binarized around its median.
    """
    # Binarize the signal to create a discrete sequence
    threshold = np.median(signal)
    binary_sequence = (signal > threshold).astype(np.uint8)

    # LZ76 Algorithm Implementation
    dictionary = set()
    current_phrase = ""
```

```

complexity = 0
for bit in binary_sequence:
    current_phrase += str(bit)
    if current_phrase not in dictionary:
        dictionary.add(current_phrase)
        current_phrase = ""
    complexity += 1

```

```

# Handle the last phrase if it's not empty
if current_phrase != "":
    complexity += 1

```

```

return complexity / len(binary_sequence)

```

```

def calculate_spectral_entropy(signal, sr, n_fft=2048, hop_length=512):

```

```

    """

```

```

    Calculates the mean spectral entropy across all frames of a signal.
    """

```

```


```

```

    # Compute the power spectrogram

```

```

    S = np.abs(librosa.stft(signal, n_fft=n_fft, hop_length=hop_length))**2

```

```

    # Add a small epsilon to avoid log(0)

```

```

    S += 1e-8

```

```

    # Normalize each frame's spectrum to a PMF

```

```

    P = S / np.sum(S, axis=0, keepdims=True)

```

```

    # Compute entropy for each frame

```

```

    entropy_per_frame = -np.sum(P * np.log2(P), axis=0)

```

```

    # Normalize by the maximum possible entropy

```

```

    max_entropy = np.log2(S.shape)

```

```

    normalized_entropy = entropy_per_frame / max_entropy

```

```

    return np.mean(normalized_entropy)

```

```

def calculate_all_features(signal, sr, n_fft=2048, hop_length=512):

```

```

    """

```

```

    Extracts a comprehensive vector of complexity features from a signal.
    """

```

```


```



```

# Time-domain features
lz_complexity = calculate_lz_complexity(signal)
rms_energy = np.mean(librosa.feature.rms(y=signal, frame_length=n_fft,
hop_length=hop_length))
zcr = np.mean(librosa.feature.zero_crossing_rate(y=signal, frame_length=n_fft,
hop_length=hop_length))

# Frequency-domain features
S = np.abs(librosa.stft(signal, n_fft=n_fft, hop_length=hop_length))
spectral_entropy = calculate_spectral_entropy(signal, sr, n_fft=n_fft,
hop_length=hop_length)
spectral_flatness = np.mean(librosa.feature.spectral_flatness(S=S))
spectral_contrast = np.mean(librosa.feature.spectral_contrast(S=S, sr=sr))

# Concatenate all features into a single vector
feature_vector = np.array([
    lz_complexity, rms_energy, zcr,
    spectral_entropy, spectral_flatness, spectral_contrast
])

return feature_vector

```

- `calculate_lz_complexity`: This function provides a custom implementation of the LZ76 algorithm. It first binarizes the continuous signal into a sequence of 0s and 1s by comparing each sample to the median value, a robust thresholding method.<sup>19</sup> It then iterates through the binary sequence, building a dictionary (implemented as a Python set for efficient lookups) of seen phrases. The complexity counter is incremented each time a new, unseen phrase is added to the dictionary. Finally, the raw complexity count is normalized by the sequence length to make it comparable across signals of different durations.
- `calculate_spectral_entropy`: This function computes spectral entropy as described in Part II. It leverages `librosa.stft` to get the spectrogram, then uses numpy for the mathematical operations: normalization to a PMF, element-wise multiplication and logarithm for the entropy calculation, and normalization by the maximum possible entropy.<sup>24</sup> The use of `np.mean` at the end aggregates the per-frame entropy values into a single feature for the entire signal.
- `calculate_all_features`: This wrapper function demonstrates the efficiency of the librosa library. Features like RMS energy, ZCR, spectral flatness, and spectral

contrast are computed with single, high-level function calls.<sup>29</sup> This abstraction is powerful; the developer can focus on the "what" (extracting spectral contrast) rather than the "how" (the complex underlying algorithm of band-splitting, sorting, and peak/valley detection).<sup>7</sup> This dramatically simplifies the code and reduces the potential for implementation errors. The function concludes by concatenating the mean of each feature across all frames into a final, fixed-size feature vector ready for the classifier.

### 3.4 Module classifier.py and main.py

These modules orchestrate the machine learning aspect of the project. classifier.py would define a class to handle model training and prediction, while main.py would use this class to execute the full pipeline.

#### Conceptual Code for main.py:

Python

```
# In main.py
from data_loader import load_audio
from feature_extractor import calculate_all_features
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
import os
import numpy as np

# 1. Load data and labels
#... code to iterate through 'data/' directory and build X (features) and y (labels)...
# for file in audio_files:
#     signal, sr = load_audio(file)
#     features = calculate_all_features(signal, sr)
#     X.append(features)
#     y.append(label_from_filename(file))
```

```

# 2. Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# 3. Scale features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# 4. Train the classifier
model = SVC(kernel='rbf', C=1.0, gamma='auto')
model.fit(X_train_scaled, y_train)

# 5. Evaluate the classifier
y_pred = model.predict(X_test_scaled)
accuracy = accuracy_score(y_test, y_pred)
print(f"Classifier Accuracy: {accuracy * 100:.2f}%")

```

This script follows standard machine learning best practices provided by scikit-learn<sup>45</sup>.

- Data is loaded and features are extracted for a labeled dataset.
- `train_test_split` separates the data to prevent the model from being evaluated on data it has already seen.
- `StandardScaler` is used to normalize the features to have zero mean and unit variance. This is crucial for distance-based algorithms like SVMs, as features with vastly different scales (e.g., RMS energy vs. spectral flatness) can otherwise disproportionately influence the model.
- An SVC (Support Vector Classifier) with an RBF kernel is instantiated, trained on the scaled training data, and then used to make predictions on the unseen test data.
- Finally, the model's performance is quantified using an appropriate metric like `accuracy_score`.

## Part IV: Strategic Applications and Future Trajectories

This final section synthesizes the analysis into a forward-looking perspective, providing actionable recommendations for improvement and exploring the project's

broader impact. It outlines how the classifier's unique approach to audio analysis can be leveraged in various real-world scenarios and how the project itself can evolve.

## 4.1 Potential Application Domains

The classifier's foundation in fundamental information-theoretic principles, rather than task-specific heuristics, makes it applicable to a wide range of domains where the underlying structure of a signal is more important than its specific semantic content.

- **Audio Forensics and Deepfake Detection:** The rise of sophisticated audio synthesis and voice conversion technologies presents a significant challenge for authentication. AI-generated speech or "deepfakes" may be perceptually convincing but often exhibit subtle statistical artifacts. For instance, they might lack the natural micro-variations of human speech, leading to an unnaturally consistent or low-complexity signature. A classifier trained to detect these subtle deviations in complexity could serve as a powerful tool for differentiating authentic recordings from synthetic or manipulated ones.<sup>12</sup>
- **Environmental Sound Classification (ESC):** This field involves identifying the type of environment from its ambient sound (e.g., a busy street, a quiet forest, an office). Different environments have distinct complexity profiles. An urban soundscape is a mixture of many independent sources, resulting in high complexity. A natural soundscape might be dominated by noise-like sounds (wind, rain) with high statistical entropy, punctuated by structured sounds (birdsong) with lower entropy but high spectral contrast.<sup>5</sup> This classifier could provide a robust method for characterizing and categorizing such soundscapes.
- **Medical Diagnostics:** Bioacoustic analysis is a growing field where sound provides a non-invasive window into physiological processes. Changes in the complexity of bodily sounds can be powerful diagnostic indicators. For example, certain cardiac conditions alter the fractal complexity of heartbeats. Pathological coughs or breathing patterns can be distinguished from healthy ones by their acoustic structure. This classifier could be adapted to analyze such signals, potentially identifying markers for respiratory or cardiovascular diseases.<sup>21</sup>
- **Music Information Retrieval (MIR):**
  - **Genre Classification:** Different musical genres often employ distinct structural and textural complexity. The dense, improvisational nature of jazz results in a different complexity signature than the repetitive, patterned

structure of minimalist electronic music. The classifier could leverage these differences for genre identification.<sup>7</sup>

- **Structural Segmentation:** Within a single piece of music, complexity can vary significantly. A sparse verse might have low complexity, while a dense, multi-layered chorus has high complexity. By tracking the evolution of the complexity feature vector over time, the system could automatically segment a song into its constituent parts (verse, chorus, bridge, solo), a key task in music analysis.<sup>31</sup>

## 4.2 Proposed Improvements and Extensions

While the current project provides a solid foundation, several strategic extensions could significantly enhance its performance, robustness, and applicability.

### 4.2.1 Integrating Perceptual Models

A primary limitation of the current feature set is that it is purely based on the signal's mathematical properties. Human perception of sound, however, is a highly non-linear and specialized process.<sup>1</sup> A signal that is algorithmically complex (e.g., white noise) can be perceptually simple (a uniform "hiss"). Bridging this "perceptual gap" is a critical next step.

This can be achieved by augmenting the feature set with psychoacoustically-motivated measures that model key aspects of the human auditory system.<sup>49</sup> For example, instead of calculating spectral features on a linear frequency scale, they could be calculated on a perceptual scale like the Mel or Bark scale. These scales better reflect how humans perceive pitch.<sup>7</sup> Furthermore, models of auditory masking—where a loud sound makes a quieter sound in its temporal or spectral vicinity inaudible—could be used to weigh the importance of different signal components, focusing the analysis on what is perceptually relevant.<sup>49</sup> This would shift the project's definition of "complexity" to be more aligned with

*perceived complexity.*

#### 4.2.2 Multiscale Complexity Analysis

The current analysis uses a fixed frame size, meaning it captures complexity at only one temporal scale. However, many important signals, especially speech and music, possess a hierarchical structure with meaningful patterns occurring at multiple time scales—from short phonemes to longer words, phrases, and musical motifs.<sup>52</sup>

A powerful extension would be to implement multiscale versions of the core complexity features. **Multi-Scale Entropy (MSE)** is a technique that involves creating "coarse-grained" versions of the time series by averaging non-overlapping blocks of samples and then calculating the sample entropy for each resulting scale.<sup>52</sup> Similarly,

**Multi-Scale Lempel-Ziv Complexity (MLZC)** can be applied to capture pattern diversity across different temporal resolutions.<sup>53</sup> Adopting a multiscale approach would provide a much richer and more robust feature vector, allowing the classifier to detect both short-term textural complexity and long-term structural complexity simultaneously.

#### 4.2.3 Advanced Classifier Architectures

The project's current pipeline, which separates handcrafted feature extraction from classification, is interpretable and effective. However, the history of progress in domains like computer vision and speech recognition has consistently shown that end-to-end deep learning models, which learn features directly from the data, ultimately achieve superior performance given a sufficiently large dataset.<sup>42</sup>

A strategic evolution of the project would involve replacing the traditional ML model with a deep neural network, such as a **Convolutional Neural Network (CNN)**, that operates directly on a time-frequency representation like a Mel spectrogram.<sup>13</sup> A CNN can be viewed as a system that automatically learns a hierarchy of relevant features. The initial layers might learn to detect simple spectral shapes and textures (analogous to ZCR or spectral flux), while deeper layers learn to combine these into more complex and abstract patterns (analogous to formants or harmonic structures). This removes the burden of manual feature engineering and allows the model to discover the most

discriminative representations for the task at hand. The existing handcrafted features would then remain invaluable as a baseline and as a tool for interpreting what the deep learning model has learned.

#### **4.2.4 Real-Time Implementation**

The current implementation is designed for the offline, batch processing of audio files. To broaden its applicability to tasks like live audio monitoring or interactive systems, it would need to be adapted for real-time processing. This would involve significant engineering effort, including:

- Refactoring the code to operate on a continuous stream of audio data, likely using a buffered or sliding-window approach.
- Optimizing the feature extraction algorithms for low latency, ensuring that the processing for one frame completes before the next frame arrives.
- Leveraging libraries designed for real-time audio I/O, such as pyAudio.<sup>42</sup>
- Carefully balancing the trade-off between frame size (which affects frequency resolution) and hop size (which affects temporal resolution and computational load).

#### **4.3 Conclusion**

The 'Complexity Audio Classifier' project represents a well-founded and insightful approach to audio classification. By grounding its methodology in the principles of algorithmic information theory, it moves beyond surface-level acoustic properties to probe the deeper structural and generative rules of sound. Its use of a complementary feature set—combining algorithmic proxies like Lempel-Ziv complexity with statistical and spectral measures—provides a robust, multi-faceted signature for characterizing audio signals. The implementation, built on the mature scientific Python ecosystem, is both elegant and efficient, demonstrating best practices in modern computational research.

The analysis reveals that the project's greatest strength lies in its potential for generalization. The source-agnostic nature of its features makes it readily adaptable to a diverse array of applications, from the critical domain of audio forensics to the

exploratory fields of medical diagnostics and music information retrieval.

The path forward, however, points toward two key evolutionary trajectories. The first is the integration of perceptual models to bridge the gap between signal-based complexity and human-perceived complexity, a step that would unlock more nuanced and subjectively relevant applications. The second is the eventual adoption of end-to-end deep learning architectures to push the boundaries of classification accuracy. The current framework serves as an excellent and interpretable foundation upon which these future advancements can be built. In its present form and with its future potential, the project stands as a compelling example of how theoretical computer science can be translated into practical and powerful tools for understanding the complex world of sound.

## Works cited

1. Complexity can facilitate visual and auditory perception - PMC, accessed June 23, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC6706302/>
2. Spectral flatness - Wikipedia, accessed June 23, 2025, [https://en.wikipedia.org/wiki/Spectral\\_flatness](https://en.wikipedia.org/wiki/Spectral_flatness)
3. Spectral flatness: quantifying how tonal or noisy sound is - Applied Mathematics Consulting, accessed June 23, 2025, <https://www.johndcook.com/blog/2016/05/03/spectral-flatness/>
4. "Kolmogorov music" by Christopher Ford - YouTube, accessed June 23, 2025, <https://www.youtube.com/watch?v=Qg3XOfioapl>
5. Effects of entropy in real-world noise on speech perception in listeners with normal hearing and hearing loss - PMC - PubMed Central, accessed June 23, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC10699887/>
6. Kolmogorov Complexity and Applications - of Marcus Hutter, accessed June 23, 2025, <http://www.hutter1.net/dagstuhl/>
7. (PDF) Music type classification by spectral contrast feature - ResearchGate, accessed June 23, 2025, [https://www.researchgate.net/publication/313484983\\_Music\\_type\\_classification\\_by\\_spectral\\_contrast\\_feature](https://www.researchgate.net/publication/313484983_Music_type_classification_by_spectral_contrast_feature)
8. pyAudioProcessing: Audio Processing, Feature Extraction, and Machine Learning Modeling, accessed June 23, 2025, <https://proceedings.scipy.org/articles/majora-212e5952-017>
9. Information theory - Wikipedia, accessed June 23, 2025, [https://en.wikipedia.org/wiki/Information\\_theory](https://en.wikipedia.org/wiki/Information_theory)
10. Spectral entropy as speech features for speech recognition - ResearchGate, accessed June 23, 2025, [https://www.researchgate.net/publication/247612912\\_Spectral\\_entropy\\_as\\_speech\\_features\\_for\\_speech\\_recognition](https://www.researchgate.net/publication/247612912_Spectral_entropy_as_speech_features_for_speech_recognition)
11. Zero-crossing rate - Wikipedia, accessed June 23, 2025, [https://en.wikipedia.org/wiki/Zero-crossing\\_rate](https://en.wikipedia.org/wiki/Zero-crossing_rate)



12. An Optimal, Universal and Agnostic Decoding Method for Message Reconstruction, Bio and Technosignature Detection - arXiv, accessed June 23, 2025, <https://arxiv.org/html/2303.16045v3>
13. Algorithm for Processing Audio Signals Using Machine Learning | Visnyk NTUU KPI Seriia, accessed June 23, 2025, <https://radap.kpi.ua/radiotechnique/article/view/1955>
14. 3.11. Zero-crossing rate - Introduction to Speech Processing, accessed June 23, 2025, [https://speechprocessingbook.aalto.fi/Representations/Zero-crossing\\_rate.html](https://speechprocessingbook.aalto.fi/Representations/Zero-crossing_rate.html)
15. Kolmogorov complexity - Wikipedia, accessed June 23, 2025, [https://en.wikipedia.org/wiki/Kolmogorov\\_complexity](https://en.wikipedia.org/wiki/Kolmogorov_complexity)
16. Kolmogorov complexity is undecidable. Does that mean that you cannot create an algorithm that compresses files into the smallest sized files possible or that you cannot prove to have created such an algorithm, though the algorithm might itself exist? - Quora, accessed June 23, 2025, <https://www.quora.com/Kolmogorov-complexity-is-undecidable-Does-that-mean-that-you-cannot-create-an-algorithm-that-compresses-files-into-the-smallest-sized-files-possible-or-that-you-cannot-prove-to-have-created-such-an-algorithm-though>
17. The Kolmogorov Test: Compression by Code Generation - arXiv, accessed June 23, 2025, <https://arxiv.org/html/2503.13992v1>
18. KOLMOGOROV COMPLEXITY IN LYRICS Teppo E. Ahonen Department of Computer Science University of Helsinki - CNRS, accessed June 23, 2025, [https://projet.liris.cnrs.fr/imagine/pub/proceedings/ICME2011/HTML/Papers/Workshops/paper\\_124.pdf](https://projet.liris.cnrs.fr/imagine/pub/proceedings/ICME2011/HTML/Papers/Workshops/paper_124.pdf)
19. Lempel–Ziv complexity - Wikipedia, accessed June 23, 2025, [https://en.wikipedia.org/wiki/Lempel%E2%80%93Ziv\\_complexity](https://en.wikipedia.org/wiki/Lempel%E2%80%93Ziv_complexity)
20. When and how to use Lempel-Ziv complexity - Information Dynamics, accessed June 23, 2025, <https://information-dynamics.github.io/complexity/information/2019/06/26/lempel-ziv.html>
21. Relationship between weighted peak frequency and LZ complexity for the AE signals generated during the testing of specimens - ResearchGate, accessed June 23, 2025, [https://www.researchgate.net/figure/Relationship-between-weighted-peak-frequency-and-LZ-complexity-for-the-AE-signals\\_fig9\\_362498736](https://www.researchgate.net/figure/Relationship-between-weighted-peak-frequency-and-LZ-complexity-for-the-AE-signals_fig9_362498736)
22. (PDF) Interpretation of the Lempel-Ziv Complexity Measure in the Context of Biomedical Signal Analysis - ResearchGate, accessed June 23, 2025, [https://www.researchgate.net/publication/6723694\\_Interpretation\\_of\\_the\\_Lempel-Ziv\\_Complexity\\_Measure\\_in\\_the\\_Context\\_of\\_Biomedical\\_Signal\\_Analysis](https://www.researchgate.net/publication/6723694_Interpretation_of_the_Lempel-Ziv_Complexity_Measure_in_the_Context_of_Biomedical_Signal_Analysis)
23. (PDF) Spectral Entropy Based Feature for Robust ASR - ResearchGate, accessed June 23, 2025, [https://www.researchgate.net/publication/4087390\\_Spectral\\_Entropy\\_Based\\_Feature\\_for\\_Robust\\_ASR](https://www.researchgate.net/publication/4087390_Spectral_Entropy_Based_Feature_for_Robust_ASR)
24. Spectral entropy | OpenAE, accessed June 23, 2025,

- <https://openae.io/standards/features/latest/spectral-entropy/>
25. What is spectral entropy? - Signal Processing Stack Exchange, accessed June 23, 2025, <https://dsp.stackexchange.com/questions/23689/what-is-spectral-entropy>
  26. What is: Zero Crossing Rate - LEARN STATISTICS EASILY, accessed June 23, 2025, <https://statisticseasily.com/glossario/what-is-zero-crossing-rate/>
  27. Auditory Selectivity for Spectral Contrast in Cortical Neurons and Behavior - PMC, accessed June 23, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC6989003/>
  28. What is RMS in the audio production world? - Major Mixing, accessed June 23, 2025, <https://majormixing.com/what-is-rms-in-audio-world/>
  29. Energy and RMSE - Music Information Retrieval, accessed June 23, 2025, <https://musicinformationretrieval.com/energy.html>
  30. Spectral flatness | OpenAE, accessed June 23, 2025, <https://openae.io/standards/features/latest/spectral-flatness/>
  31. Using a Spectral Flatness Based Feature for Audio Segmentation and Retrieval, accessed June 23, 2025, <https://ismir2000.ismir.net/posters/izmirli.pdf>
  32. SpectralContrast — Essentia 2.1-beta6-dev documentation, accessed June 23, 2025, [https://essentia.upf.edu/reference/std\\_SpectralContrast.html](https://essentia.upf.edu/reference/std_SpectralContrast.html)
  33. A Tutorial on Spectral Feature Extraction for Audio Analytics -, accessed June 23, 2025, <https://analyticsindiamag.com/ai-trends/a-tutorial-on-spectral-feature-extraction-for-audio-analytics/>
  34. Overview of Spectral Contrast Feature calculation - ResearchGate, accessed June 23, 2025, [https://www.researchgate.net/figure/Overview-of-Spectral-Contrast-Feature-calculation\\_fig1\\_220723537](https://www.researchgate.net/figure/Overview-of-Spectral-Contrast-Feature-calculation_fig1_220723537)
  35. Zero-crossing rate | OpenAE, accessed June 23, 2025, <https://openae.io/standards/features/latest/zero-crossing-rate/>
  36. zerocrossrate - Zero-crossing rate - MATLAB - MathWorks, accessed June 23, 2025, <https://www.mathworks.com/help/signal/ref/zerocrossrate.html>
  37. What is an Audio Feature? - Meyda, accessed June 23, 2025, <https://meyda.js.org/audio-features.html>
  38. Add RMSEnergyExtractor [audio feature extraction]? · Issue #698 · IAHispano/Applio - GitHub, accessed June 23, 2025, <https://github.com/IAHispano/Applio/issues/698>
  39. Understanding RMS in Audio - Online Mixing and Mastering Services, accessed June 23, 2025, <https://mixandmastermysong.com/understanding-rms-in-audio/>
  40. Comparison of the RMS Energy and the Amplitude Envelope - Analytics Vidhya, accessed June 23, 2025, <https://www.analyticsvidhya.com/blog/2022/05/comparison-of-the-rms-energy-and-the-amplitude-envelope/>
  41. Hands-On Guide To Librosa For Handling Audio Files - Analytics Vidhya, accessed June 23, 2025, <https://www.analyticsvidhya.com/blog/2024/01/hands-on-guide-to-librosa-for-handling-audio-files/>

42. 10 Python Libraries for Audio Processing - CloudDevs, accessed June 23, 2025, <https://cloudd devs.com/python/libraries-for-audio-processing/>
43. librosa 0.11.0 documentation, accessed June 23, 2025, <https://librosa.org/doc/>
44. libAudioFlux/audioFlux: A library for audio and music analysis, feature extraction. - GitHub, accessed June 23, 2025, <https://github.com/libAudioFlux/audioFlux>
45. An Evaluation of Entropy Measures for Microphone Identification - MDPI, accessed June 23, 2025, <https://www.mdpi.com/1099-4300/22/11/1235>
46. Audio processing in Python with Feature Extraction for machine learning - YouTube, accessed June 23, 2025, <https://www.youtube.com/watch?v=vbhlEMcb7RQ>
47. zcr - Zero Crossing Rate - Music Information Retrieval, accessed June 23, 2025, <https://musicinformationretrieval.com/zcr.html>
48. How We Hear: The Perception and Neural Coding of Sound - PMC - PubMed Central, accessed June 23, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC5819010/>
49. Perceptual Audio Coding: A 40-Year Historical Perspective - arXiv, accessed June 23, 2025, <https://arxiv.org/html/2504.16223v1>
50. Psychoacoustic Models for Perceptual Audio Coding—A Tutorial Review - MDPI, accessed June 23, 2025, <https://www.mdpi.com/2076-3417/9/14/2854>
51. Objective Quality Assessment of Perceptually Coded Audio Signals - OPEN FAU, accessed June 23, 2025, <https://open.fau.de/handle/openfau/22761>
52. Understanding Multiscale Entropy - Sapien Labs | Shaping the Future of Mind Health, accessed June 23, 2025, <https://sapienlabs.org/understanding-multiscale-entropy/>
53. Multi-scale permutation Lempel-Ziv complexity and its application in feature extraction for Ship-radiated noise - Frontiers, accessed June 23, 2025, <https://www.frontiersin.org/journals/marine-science/articles/10.3389/fmars.2022.1047332/full>